

nkxqosbud

October 26, 2023

Naan Mudhalva Phase4 Development: Part2

Sentiment Analysis for Marketing

```
[32]: pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.34.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.12.4)
Requirement already satisfied: huggingface-hub<1.0,>=0.16.4 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.17.3)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.23.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (23.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.1)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2023.6.3)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.31.0)
Requirement already satisfied: tokenizers<0.15,>=0.14 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.14.1)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.4.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.66.1)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.16.4->transformers) (2023.6.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.16.4->transformers) (4.5.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.3.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2.0.7)
```

Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests->transformers)
(2023.7.22)

```
[33]: from PIL import Image
from sklearn import svm
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import roc_curve
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
from transformers import BertTokenizer, BertForSequenceClassification, pipeline
from sklearn.model_selection import train_test_split
from transformers import pipeline

import collections
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np
import operator
import pandas as pd
import torch

from transformers import BertTokenizer, BertForSequenceClassification, AdamW
from torch.utils.data import DataLoader, Dataset
from torch.nn import functional as F
import matplotlib.pyplot as plt
```

```
[34]: tweets = pd.read_csv('Tweets.csv')
```

```
[35]: tweets.head(5)
```

```
[35]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	\
0	570306133677760513	neutral	1.0000	
1	570301130888122368	positive	0.3486	
2	570301083672813571	neutral	0.6837	
3	570301031407624196	negative	1.0000	
4	570300817074462722	negative	1.0000	

	negativereason	negativereason_confidence	airline	\
0	NaN	NaN	Virgin America	
1	NaN	0.0000	Virgin America	
2	NaN	NaN	Virgin America	
3	Bad Flight	0.7033	Virgin America	
4	Can't Tell	1.0000	Virgin America	

	airline_sentiment_gold	name	negativereason_gold	retweet_count	\
0	NaN	cairdin	NaN	0	
1	NaN	jnardino	NaN	0	
2	NaN	yvonnalynn	NaN	0	
3	NaN	jnardino	NaN	0	
4	NaN	jnardino	NaN	0	

	text	tweet_coord	\
0	@VirginAmerica What @dhepburn said.	NaN	
1	@VirginAmerica plus you've added commercials t...	NaN	
2	@VirginAmerica I didn't today... Must mean I n...	NaN	
3	@VirginAmerica it's really aggressive to blast...	NaN	
4	@VirginAmerica and it's a really big bad thing...	NaN	

	tweet_created	tweet_location	user_timezone
0	2015-02-24 11:35:52 -0800	NaN	Eastern Time (US & Canada)
1	2015-02-24 11:15:59 -0800	NaN	Pacific Time (US & Canada)
2	2015-02-24 11:15:48 -0800	Lets Play	Central Time (US & Canada)
3	2015-02-24 11:15:36 -0800	NaN	Pacific Time (US & Canada)
4	2015-02-24 11:14:45 -0800	NaN	Pacific Time (US & Canada)

```
[36]: tweets['negativereason_gold'].value_counts()
```

```
[36]: Customer Service Issue      12
Late Flight                      4
Can't Tell                       3
Cancelled Flight                 3
Cancelled Flight\nCustomer Service Issue  2
Late Flight\nFlight Attendant Complaints  1
Late Flight\nLost Luggage          1
Bad Flight                      1
Lost Luggage\nDamaged Luggage      1
Late Flight\nCancelled Flight       1
Flight Attendant Complaints       1
Customer Service Issue\nLost Luggage  1
Customer Service Issue\nCan't Tell  1
Name: negativereason_gold, dtype: int64
```

```
[37]: tweets['airline_sentiment_gold'].value_counts()
```

```
[37]: negative      32
positive      5
neutral       3
Name: airline_sentiment_gold, dtype: int64
```

```
[38]: tweets['retweet_count'].value_counts()
```

```
[38]: 0      13873
      1       640
      2        66
      3        22
      4        17
      5         5
      7         3
      6         3
      22         2
      8         1
      32         1
      28         1
      9         1
      18         1
      11         1
      31         1
      15         1
      44         1
      Name: retweet_count, dtype: int64
```

```
[39]: tweets.drop('negativereason_gold', axis=1, inplace=True)
      tweets.drop('airline_sentiment_gold', axis=1, inplace=True)
      tweets.drop('retweet_count', axis=1, inplace=True)
      tweets.drop('tweet_coord', axis=1, inplace=True)

      tweets.drop('tweet_location', axis=1, inplace=True)
      tweets.drop('tweet_created', axis=1, inplace=True)
      tweets.drop('user_timezone', axis=1, inplace=True)
      tweets.drop('name', axis=1, inplace=True)

      list(tweets.columns)
```

```
[39]: ['tweet_id',
      'airline_sentiment',
      'airline_sentiment_confidence',
      'negativereason',
      'negativereason_confidence',
      'airline',
      'text']
```

```
[40]: unmeaningful = ['i', 'you', 'me', 'to', 'the', 'a', 'my', 'is', 'in', 'and',
      ↪ 'for', 'on', 'of',
      'your', 'so', 'was', 'have', 'it', 'at', 'with', 'that',
      ↪ 'from', 'do', 'get',
      'but', 'this', 'can', 'just', 'they', 'we', 'are', 'an', 'be',
      ↪ 'i'm', 'will',
```

```

        'if', 'had', 'our', 'about', 'there', 'has', 'been', '-', 'by',
↪ 'like', 'or',
        'as', 'he', 'she', 'it', 'us', 'has', 'i've', 'it's', 'don't',
↪ 'would', 'am',
        'flight', 'customer', 'any', 'very', 'didn't', 'you've',
↪ 'thing', 'take',
        'other', 'u', '', ' ' ]

```

```

[41]: def clean_text(str_in):

    res = ""
    str_in = str_in.lower()
    str_arr = str_in.split(' ')
    for word in str_arr:

        word = word.lower()

        if '@' in word or word == '' or word[:1] == '&':
            continue
        if word.lower() in unmeaningful:
            continue
        if word.isnumeric():
            continue
        res = res + " " + word
    return res

```

```

[42]: tweets["text"] = tweets["text"].apply(clean_text)

```

```

[43]: tweets.head(5)

```

```

[43]:      tweet_id  airline_sentiment  airline_sentiment_confidence  \
0  570306133677760513             neutral                1.0000
1  570301130888122368             positive                0.3486
2  570301083672813571             neutral                0.6837
3  570301031407624196             negative                1.0000
4  570300817074462722             negative                1.0000

```

```

      negativereason  negativereason_confidence      airline  \
0              NaN              NaN  Virgin America
1              NaN              0.0000  Virgin America
2              NaN              NaN  Virgin America
3      Bad Flight              0.7033  Virgin America
4      Can't Tell              1.0000  Virgin America

```

```

                                text
0                                what said.
1      plus added commercials experience... tacky.

```

```

2             today... must mean need another trip!
3  really aggressive blast obnoxious "entertainm...
4                                     really big bad

```

```

[57]: data = tweets
data['airline_sentiment'] = data['airline_sentiment'].astype('category')

X = data['text']
y = data['airline_sentiment']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

X_train_encodings = tokenizer(list(X_train), padding=True, truncation=True,
↳return_tensors='pt', max_length=128)
X_test_encodings = tokenizer(list(X_test), padding=True, truncation=True,
↳return_tensors='pt', max_length=128)
y_train_encodings = torch.tensor(y_train.cat.codes.values)
y_test_encodings = torch.tensor(y_test.cat.codes.values)

```

```

[47]: class SentimentDataset(Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = labels

    def __getitem__(self, idx):
        item = {key: val[idx] for key, val in self.encodings.items()}
        item['labels'] = self.labels[idx]
        return item

    def __len__(self):
        return len(self.labels)

```

```

[48]: train_dataset = SentimentDataset(X_train_encodings, y_train_encodings)
test_dataset = SentimentDataset(X_test_encodings, y_test_encodings)

```

```

[49]: model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
↳num_labels=3)
optimizer = AdamW(model.parameters(), lr=1e-5)

```

Downloading model.safetensors: 0%| | 0.00/440M [00:00<?, ?B/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

```
['classifier.bias', 'classifier.weight']
```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411:
```

FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Use the PyTorch implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True` to disable this warning

```
warnings.warn(
```

```
[50]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
model.train()
```

```
[50]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
)
```

```

    )
    )
    )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.1, inplace=False)
  (classifier): Linear(in_features=768, out_features=3, bias=True)
)

```

```

[53]: model.eval()
test_loader = DataLoader(test_dataset, batch_size=64)
predictions = []

```

```

[54]: for batch in test_loader:
    input_ids = batch['input_ids'].to(device)
    attention_mask = batch['attention_mask'].to(device)
    with torch.no_grad():
        outputs = model(input_ids, attention_mask=attention_mask)
    logits = outputs.logits
    predicted_labels = F.softmax(logits, dim=1).argmax(dim=1)
    predictions.extend(predicted_labels.cpu().numpy())

```

```

[55]: accuracy = accuracy_score(y_test_encodings, predictions)
print("Accuracy:", accuracy)
print(classification_report(y_test_encodings, predictions))

```

Accuracy: 0.6441256830601093

	precision	recall	f1-score	support
0	0.64	1.00	0.78	1889
1	0.00	0.00	0.00	580
2	0.00	0.00	0.00	459
accuracy			0.64	2928
macro avg	0.21	0.33	0.26	2928
weighted avg	0.42	0.64	0.51	2928

```

[56]: sentiment_counts = data['airline_sentiment'].value_counts()
plt.figure(figsize=(8, 5))
plt.bar(sentiment_counts.index, sentiment_counts.values, color=['red', 'green', 'blue'])
plt.title('Sentiment Distribution')

```



```
plt.xlabel('Sentiment')  
plt.ylabel('Count')  
plt.show()
```

